

Hw7

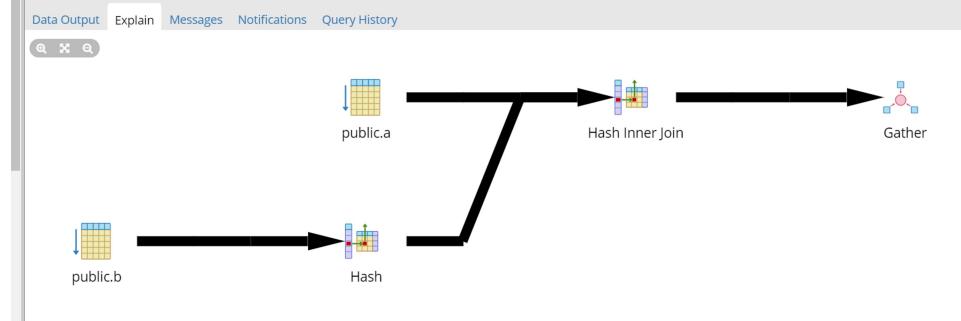
Sunday, November 4, 2018 16:19

I used Postgresql for this homework. I've included the jdbc code I used to generate the test data, and the SQL commands I used to generated the data, copy it to tables A,B,C,A',B', and C' and to create the secondary indexes.

Section 1: Explain Query Plans

1. A.pk = B.pk

```
22 CREATE INDEX c_ten_index ON t_prime(ten);  
23  
24 /* SECTION 1 */  
25 /* */  
26 SELECT * FROM A JOIN B ON A.pk = B.pk;  
27
```



System Chooses Hash Join. Hashes B and scans over A to do hash inner join

Expected Time:1338135.43

Expected Size:5000016 rows

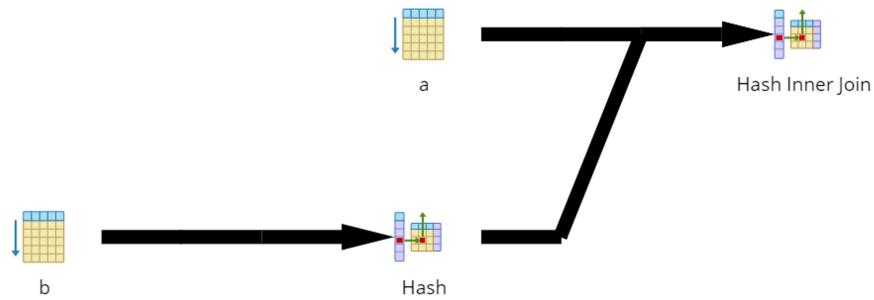
T(A) = 5000000 V(A,pk) = 5000000

T(B) = 5000000 V(B,pk) = 5000000

The selected plan is hash join.

2. A.ht = B.ht

```
/*2*/  
SELECT * FROM A JOIN B ON A.ht = B.ht;  
/*3*/
```



$$T(A) = 5000000 \quad V(A,ht) = 100000$$

$$T(B) = 5000000 \quad V(B,ht) = 100000$$

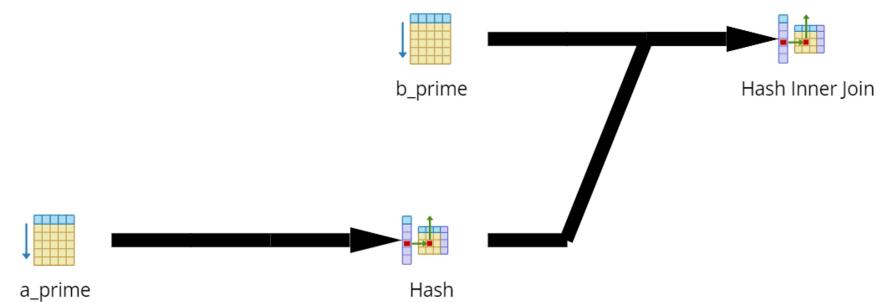
3. $A'.ht = B'.ht$

```

31 /*3*/
32 SELECT * FROM A_prime JOIN B_prime ON A_prime.ht =B_prime.ht;
33
34 /*4*/
35 SELECT * FROM A JOIN B ON A.ten =B.ten;

```

Data Output Explain Messages Notifications Query History



$$T(A') = 5000000 \quad V(A',ht) = 100000$$

$$T(B') = 5000000 \quad V(B',ht) = 100000$$

$V(R,attrib)$ did not change the query plan across 1,2,3. The secondary index has no impact on the selected query plan.

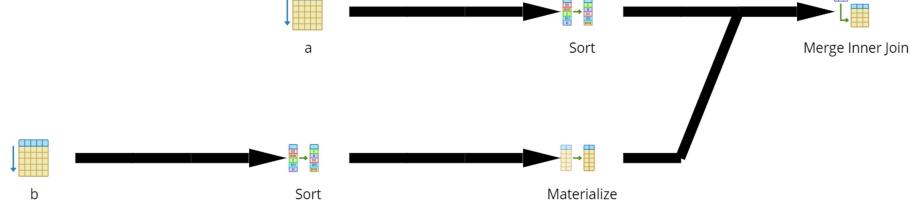
4. $A.ten = B.ten$

```

33 /*4*/
34 SELECT * FROM A JOIN B ON A.ten =B.ten;
35
36 /*5*/
37

```

Data Output Explain Messages Notifications Query History



$$T(A) = 5000000 \quad V(A,ten) = 10$$

$$T(B) = 5000000 \quad V(B,ten) = 10$$

Query plan selects Merge join

5. $A'.ten = B'.ten$

```

38 /*5*/
39 SELECT * FROM A_prime JOIN B_prime ON A_prime.ten =B_prime.ten;
40
41 /*6*/

```

Data Output Explain Messages Notifications Query History



```

36
37 /*5*/
38 SELECT * FROM A_prime JOIN B_prime ON A_prime.ten =B_prime.ten;
39
40 /*6*/

```

Data Output Explain Messages Notifications Query History

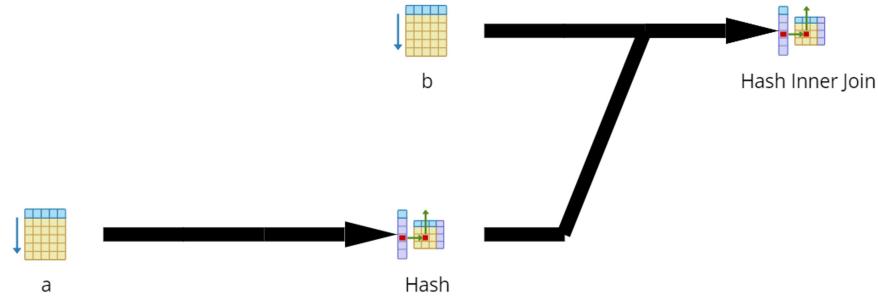


```

42
43 /*7*/
44 SELECT * FROM B JOIN A ON B.ten =A.ht;
45
46 /*8*/

```

Data Output Explain Messages Notifications Query History



$$T(A) = 5000000 \quad V(A, ht) = 100000$$

$$T(B) = 5000000 \quad V(B, ten) = 10$$

Selected plan is hash join again. The optimizer chose to hash a, then join b regardless of the order that a,b appear in the JOIN statement

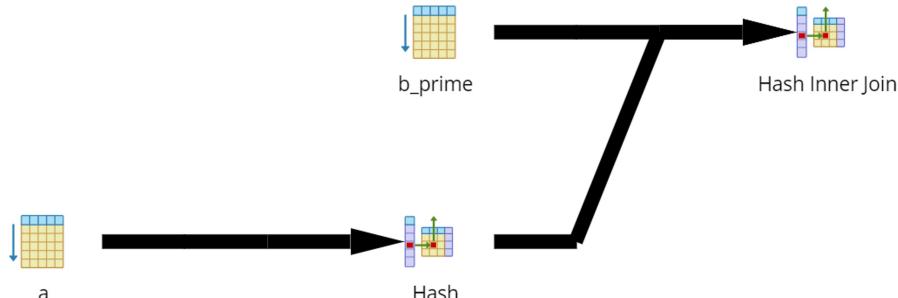
8. A.ht = B'.ten

```

/*8*/
SELECT * FROM A JOIN B_prime ON A.ht =B_prime.ten;
/*9*/

```

Data Output Explain Messages Notifications Query History



$$T(A) = 5000000 \quad V(A, ht) = 100000$$

$$T(B') = 5000000 \quad V(B', ten) = 10$$

Selected plan is hash join

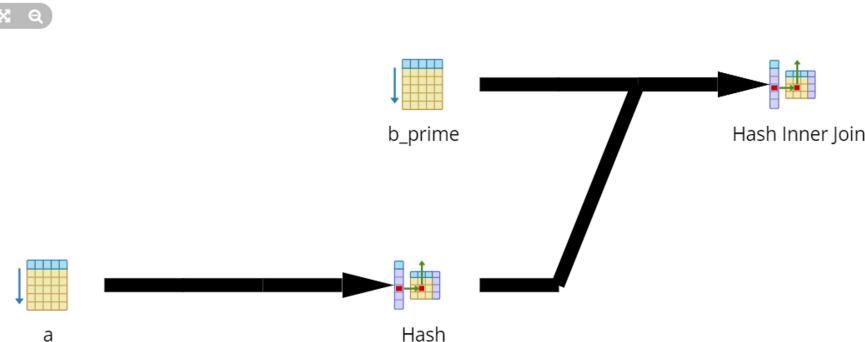
9. B'.ten = A.ht

```

47
48   SELECT * FROM A JOIN B_prime ON A.ht = B_prime.ten;
49   /*9*/
50   SELECT * FROM B_prime JOIN A ON B_prime.ten = A.ht;
51
52   /* 3 way joins */
53   /*10*/
54   SELECT * FROM A JOIN B ON A.ht = B.ht JOIN C ON B.ht = C.ht;

```

Data Output Explain Messages Notifications Query History



$$T(A) = 5000000 \quad V(A,ht) = 100000$$

$$T(B') = 5000000 \quad V(B',ten) = 10$$

Selected plan is hash join. The optimizer again chose its own order regardless of how a,b appear in the JOIN statement

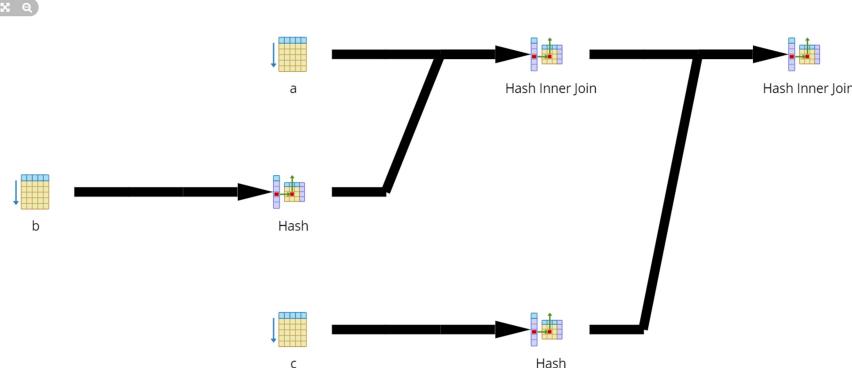
10. $A.ht = B.ht = C.ht$

```

51
52   /* 3 way joins */
53   /*10*/
54   SELECT * FROM A JOIN B ON A.ht = B.ht JOIN C ON B.ht = C.ht;
55
56   /*11*/
57   SELECT * FROM A JOIN B ON A.ten = B.ten JOIN C ON B.ten = C.ten;

```

Data Output Explain Messages Notifications Query History



$$T(A) = 5000000 \quad V(A,ht) = 100000$$

$$T(B) = 5000000 \quad V(B,ht) = 100000$$

$$T(c) = 5000000 \quad V(C,ht) = 100000$$

Selected plan: hash join A,b, hash join c

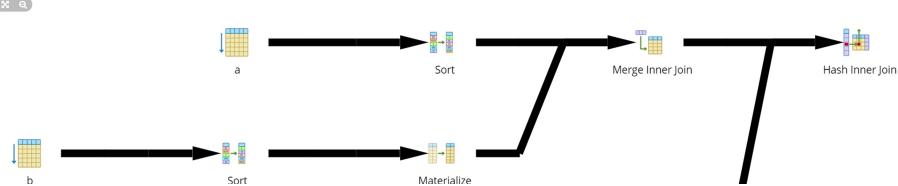
11. $A.ten = B.ten = C.ten$

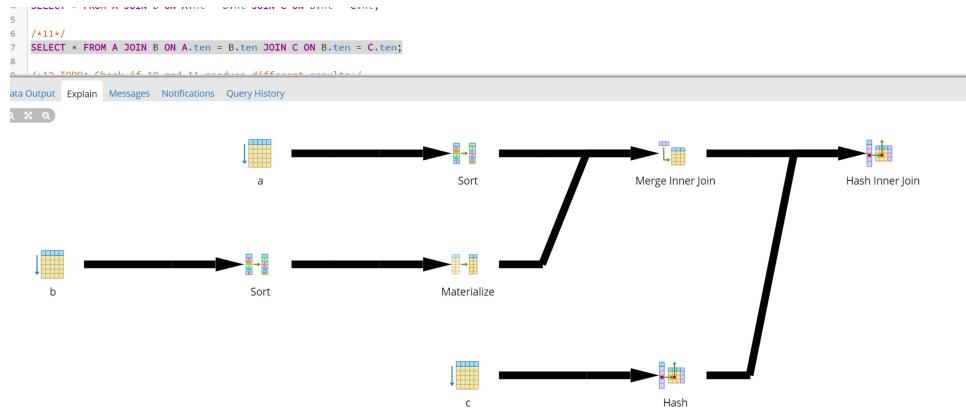
```

5
6   /*11*/
7   SELECT * FROM A JOIN B ON A.ten = B.ten JOIN C ON B.ten = C.ten;
8
9   /*12. YODA: Oracle 24.10 and 19 produce different answers!*/

```

Data Output Explain Messages Notifications Query History





$$T(A) = 5000000 \quad V(A,ht) = 10$$

$$T(B) = 5000000 \quad V(B,ht) = 10$$

$$T(C) = 5000000 \quad V(C,ht) = 10$$

Selected plan: merge join A,B, hash join c

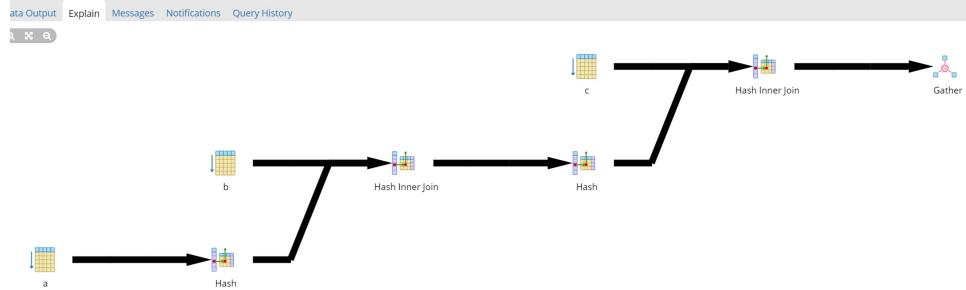
12. Do they produce different plans?

Yes, the plans are different. I found the plan changes on the ot attribute when $V(R,ot) = 1000$

$$A.ot = B.ot = C.ot$$

13. $A.pk = B.ht = C.ht$

```
/*12*/
SELECT * FROM A JOIN B ON A.ot = B.ot JOIN C ON B.ot = C.ot; */
/*13*/
SELECT * FROM A JOIN B ON A.pk = B.ht JOIN C ON B.ht = C.ht
/*14*/
SELECT * FROM A JOIN B ON A.pk = B.ht JOIN C ON B.ht = C.ht
```



Selected plan is to hash join a,b then hash join c

$$T(A) = 5000000 \quad V(A,pk) = 5000000$$

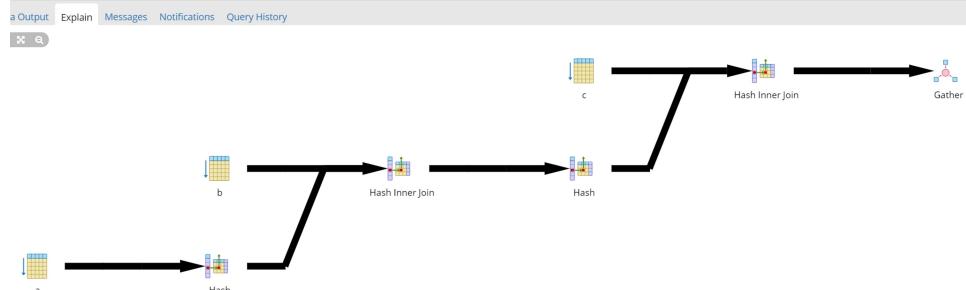
$$T(B) = 5000000 \quad V(B,ht) = 1000000$$

$$T(C) = 5000000 \quad V(C,ht) = 1000000$$

14. $A.pk = B.hund = C.hund$

```
/*14*/
SELECT * FROM A JOIN B ON A.pk = B.hund JOIN C ON B.hund = C.hund
/*15 TODO: Check if 13 and 14 produce different results */
/*16*/
SELECT * FROM A_prime JOIN B_prime ON A_prime.ht = B_prime.ht JOIN C_prime ON B_prime.ht = C_prime.ht
/*17*/

```



Selected plan is to hash join a,b then hash join c

$$T(A) = 5000000 \quad V(A,pk) = 5000000$$

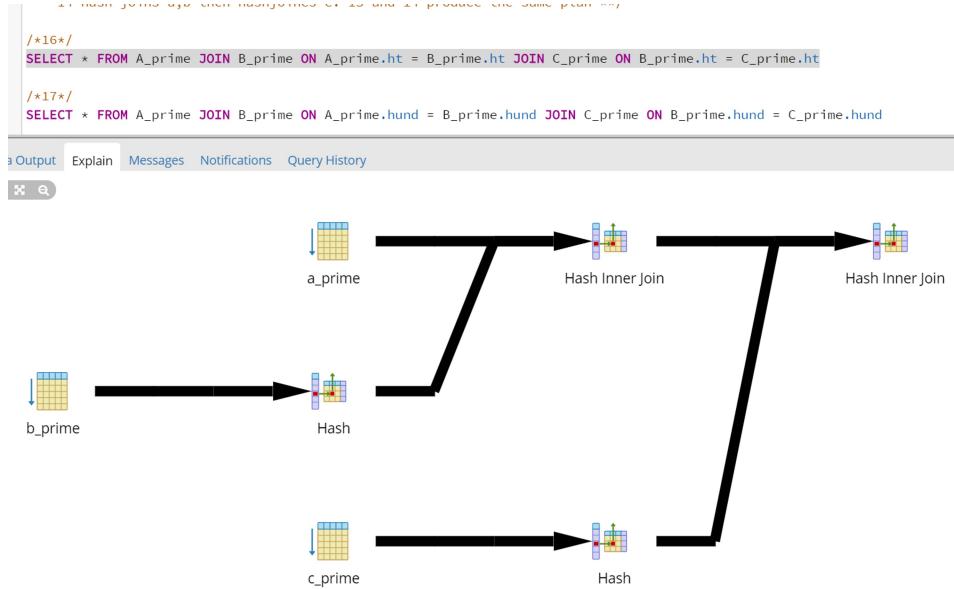
$$T(B) = 5000000 \quad V(B,hund) = 100$$

$$T(C) = 5000000 \quad V(C,hund) = 100$$

15. Do they produce different plans?

No, 13 and 14 produce the same plan.

16. $A'.ht = B'.ht = C'.ht$



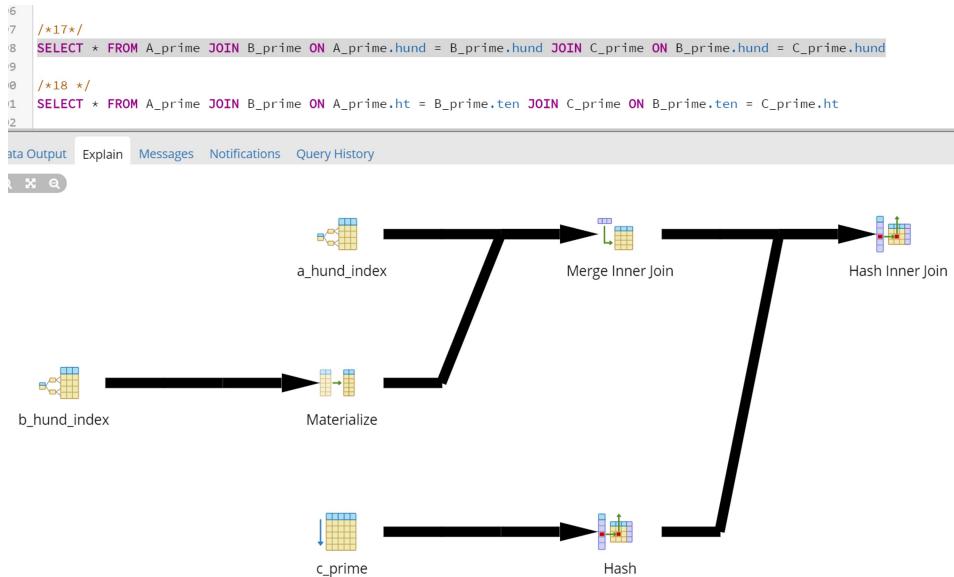
$$T(A) = 5000000 \quad V(A,ht) = 100000$$

$$T(B) = 5000000 \quad V(B,ht) = 100000$$

$$T(C) = 5000000 \quad V(C,ht) = 100000$$

Selected plan: hash join A,b, hash join c. The query plan does not change from 10, so the secondary index has no effect on the plan

17. $A'.hund = B'.hund = C'.hund$



$$T(A) = 5000000 \quad V(A,hund) = 100$$

$$T(B) = 5000000 \quad V(B,hund) = 100$$

$$T(C) = 5000000 \quad V(C,hund) = 100$$

Selected plan: merge join A,b, hash join c. The query plan changes from 16.

18. Do 16 and 17 produce different results?

Yes. 16 chooses hash join to join a,b. 17 chooses merge join.

Section 2:

19. $A'.ht = C'.ot \text{ AND}$

$A.ht = A'.ten \text{ AND}$

$B.pk = C'.hund \text{ AND}$

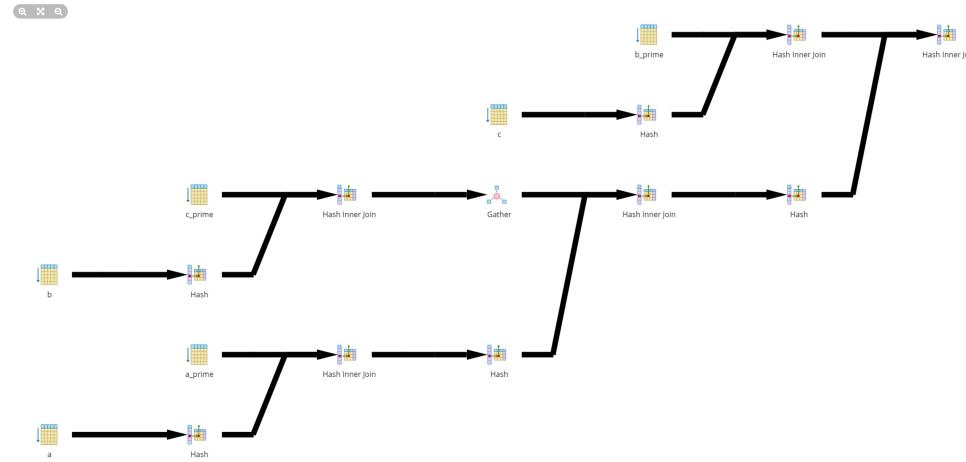
$A'.ht = C'.ot \text{ AND}$

$B'.ten = C.ot \text{ AND}$

$B'.hund = A.ht$

```
182
183 /*+19*/
184 SELECT * FROM A_prime,C_prime,A,B,B_prime,C
185 WHERE
186 A_prime.ht = C_prime.ot AND
187 A.ht = A_prime.ten AND
188 B_pk = C_prime.hund AND
189 A_prime.ht = C_prime.ot AND
190 B_prime.ten = C.ot AND
191 B_prime.hund = A.ht
192
```

Data Output Explain Messages Notifications Query History



The estimated size in the last hash inner join is over 3 quadrillion rows (3,172,245,629,574,755)

20. $A'.ht = C'.ot \text{ AND}$

$A'.ten = 5 \text{ AND}$

$A.ht < A'.ten \text{ AND}$

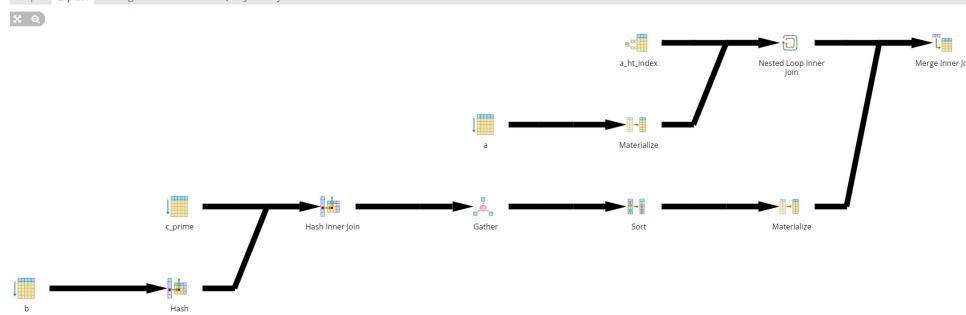
$B.pk = C'.hund \text{ AND}$

$B.ot < 500 \text{ AND}$

$A'.ht = C'.ot$

```
/*+28*/
SELECT * FROM A_prime,C_prime,A,B
WHERE
A_prime.ht = C_prime.ot AND
A_prime.ten = 5 AND
A.ht < A_prime.ten AND
B_pk = C_prime.hund AND
B.ot < 500 AND
A_prime.ht = C_prime.ot
```

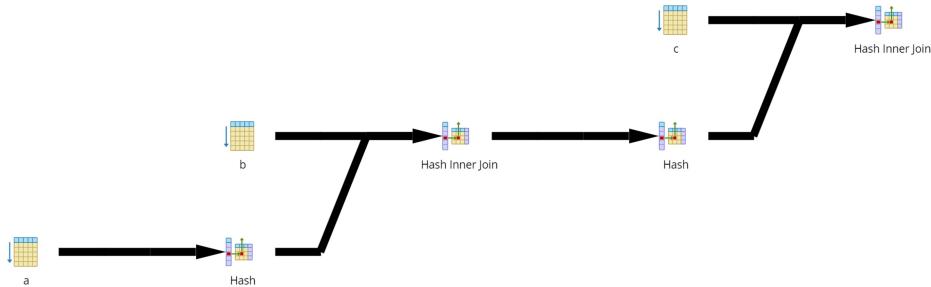
Data Output Explain Messages Notifications Query History



The estimated size of the last merge inner join is: over 21 trillion rows(21,554,197,796,542)!
The compiler did not notice the equivalent predicate.

21. A.ht = B.tt = C.ot

```
14 //+21*/  
15 SELECT * FROM A  
16 JOIN B ON A.h_t = B.b_tt  
17 JOIN C ON B.c_ott = C.o_tt  
18 //+22*/  
19 SELECT * FROM C  
20 JOIN B ON C.o_tt = B.b_tt
```



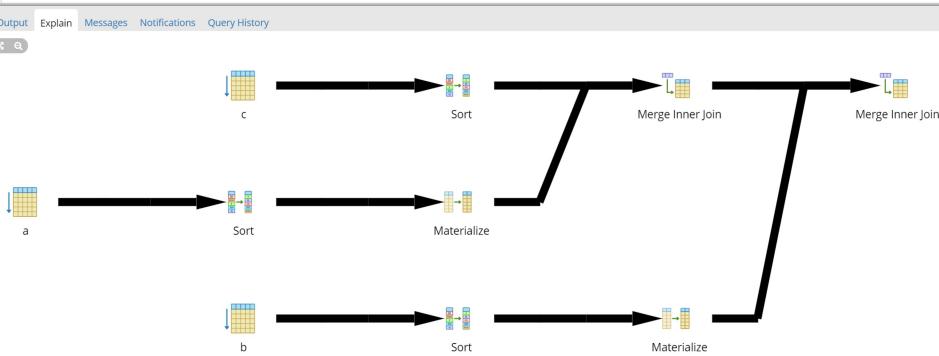
$$T(A) = 5000000 \quad V(A, ht) = 100000$$

$$T(B) = 5000000 \quad V(B,tt) = 10000$$

$$T(C) = 5000000 \quad V(C, ot) = 1000$$

$$22. C_{ot} = B_{tt} = A_{ht}$$

```
17  
18 /*22*/  
19 SELECT * FROM C  
20 JOIN B ON C.ot = B.tt  
21 JOIN A ON B.tt = A.ht  
22
```



The query plan changes between 21 and 22 even though the only difference is the order of the tables in the query.

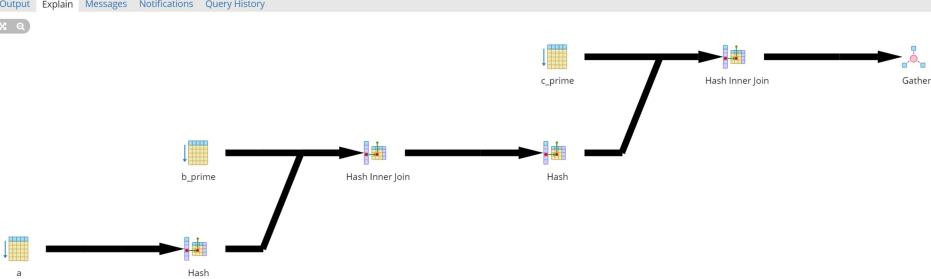
$$T(A) = 5000000 \quad V(A,ht) = 100000$$

$$T(B) = 5000000 \quad V(B,tt) = 10000$$

$$T(C) = 5000000 \quad V(C, ot) = 1000$$

23. A.pk = B'.tt = C'.ot

```
32
33  /*+3*/+
34  SELECT * FROM A
35  JOIN B_prime ON A.pk = B_prime.tt
36  JOIN C_prime ON B_prime.tt = C_prime.ot
37
38  /*+4*/
39  SELECT *FROM C_prime
40  JOIN B_prime ON C_prime.ot = B_prime.tt
41  JOIN A ON B_prime.tt = A.pk
42
```

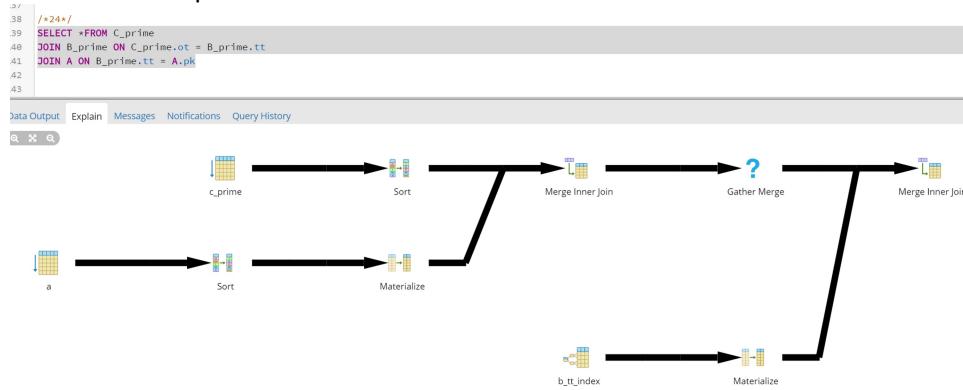


$T(A) = 5000000$ $T(A,pk) = 5000000$

$T(B) = 5000000$ $T(B,tt) = 10000$

$T(C) = 5000000$ $T(C,ot) = 1000$

24. $C'.ot = B'.tt = A.pk$



$T(A) = 5000000$ $T(A,pk) = 5000000$

$T(B) = 5000000$ $T(B,tt) = 10000$

$T(C) = 5000000$ $T(C,ot) = 1000$

Again, the query plan changes between 23 and 24 even though the only difference is the order of the tables in the query