

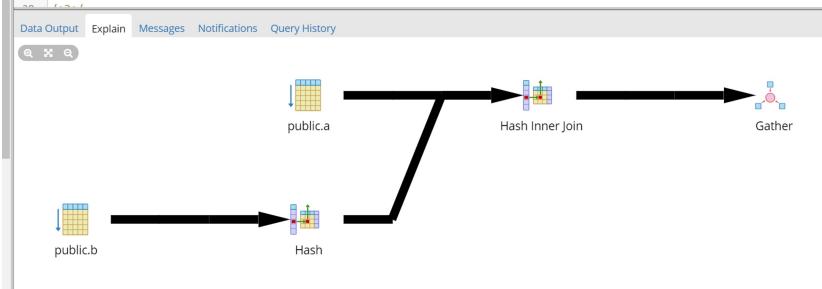
Appendix

Sunday, November 4, 2018 16:19

Section 1: Explain Query Plans

1. A.pk = B.pk

```
22 CREATE INDEX c_ten_index ON c_prime(ten);
23
24 /** SECTION 1 ***/
25 /*1*/
26 SELECT * FROM A JOIN B ON A.pk =B.pk;
27
```



System Chooses Hash Join. Hashes B and scans over A to do hash inner join

Expected Time:1338135.43

Expected Size:5000016 rows

T(A) = 5000000 V(A,pk) = 5000000

T(B) = 5000000 V(B,pk) = 5000000

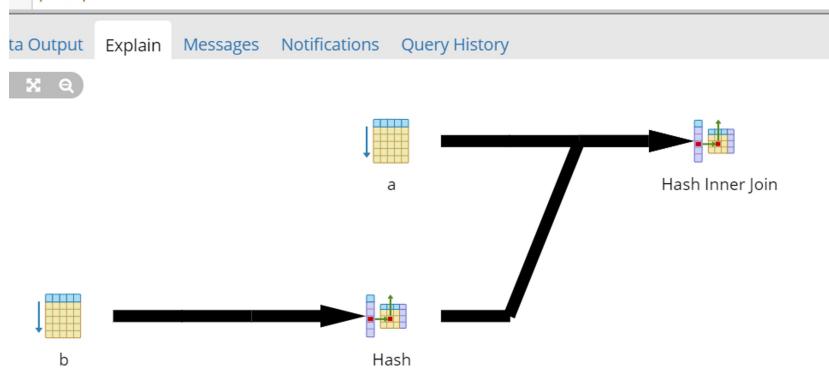
The selected plan is hash join.

2. A.ht = B.ht

```
/*2*/
SELECT * FROM A JOIN B ON A.ht =B.ht;

/*3*/

```



T(A) = 5000000 V(A,ht) = 100000
T(B) = 5000000 V(B,ht) = 100000

3. A'.ht = B'.ht

```
/*4*/
SELECT * FROM A_prime JOIN B_prime ON A_prime.ht =B_prime.ht;
/*5*/
SELECT * FROM A JOIN B ON A.ten =B.ten;
```

Some plan even though A,B have indirection ht.



```

1 /*3*/
2 SELECT * FROM A_prime JOIN B_prime ON A_prime.ht = B_prime.ht;
3
4 /*4*/
5 SELECT * FROM A JOIN B ON A.ten = B.ten;

```

Data Output Explain Messages Notifications Query History



Search

Reset

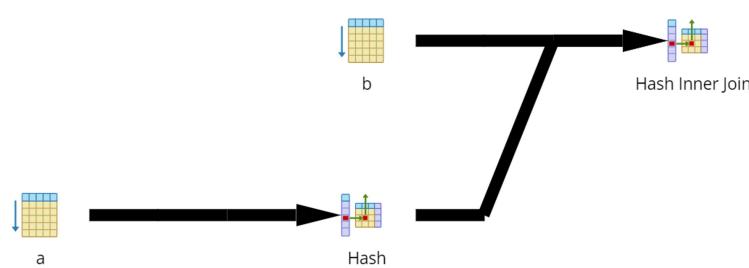
Search

```

39
40 /*6*/
41 SELECT * FROM A JOIN B ON A.ht =B.ten;
42
43 /*7*/

```

Data Output Explain Messages Notifications Query History



$T(A) = 5000000$ $V(A,ht) = 100000$

$T(B) = 5000000$ $V(B,ten) = 10$

Selected plan is hash join

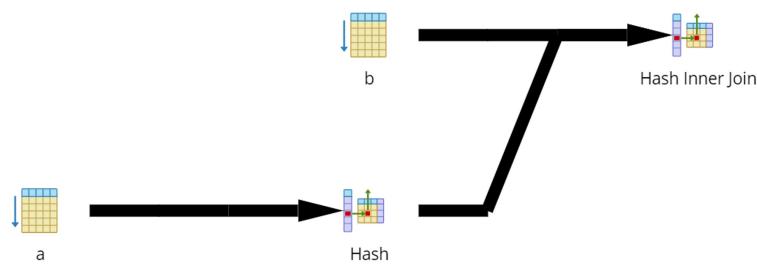
7. $B.ten = A.ht$

```

42
43 /*7*/
44 SELECT * FROM B JOIN A ON B.ten =A.ht;
45
46 /*8*/

```

Data Output Explain Messages Notifications Query History



$T(A) = 5000000$ $V(A,ht) = 100000$

$T(B) = 5000000$ $V(B,ten) = 10$

Selected plan is hash join again. The optimizer chose to hash a, then join b regardless of the order that a,b appear in the JOIN statement

8. $A.ht = B.ten$

```
/*8*/
SELECT * FROM A JOIN B_prime ON A.ht = B_prime.ten;
```

```
/*9*/
```

Data Output Explain Messages Notifications Query History



b_prime



b_prime



Hash Inner Join



a



Hash



Hash Inner Join

$T(A) = 5000000 V(A,ht) = 100000$

$T(B') = 5000000 V(B',ten) = 10$

Selected plan is hash join

9. $B'.ten = A.ht$

```
/*7*/
SELECT * FROM A JOIN B_prime ON A.ht = B_prime.ten;
```

```
48
49 /*9*/
50 SELECT * FROM B_prime JOIN A ON B_prime.ten = A.ht;
51
52 /* 3 way joins */
53 /*10*/
54 SELECT * FROM A JOIN B ON A.ht = B.ten JOIN C ON B.ten = C.ht;
```

Data Output Explain Messages Notifications Query History



b_prime



b_prime



Hash Inner Join



a



Hash



Hash Inner Join

$T(A) = 5000000 V(A,ht) = 100000$

$T(B') = 5000000 V(B',ten) = 10$

Selected plan is hash join. The optimizer again chose its own order regardless of how a,b appear in the JOIN statement

10. $A.ht = B.ht = C.ht$

```
12 /* 3 way joins */
13 /*10*/
14 SELECT * FROM A JOIN B ON A.ht = B.ht JOIN C ON B.ht = C.ht;
15
16 /*11*/
17 SELECT * FROM A JOIN B ON A.ten = B.ten JOIN C ON B.ten = C.ten;
```

Data Output Explain Messages Notifications Query History



a



a



Hash Inner Join



Hash Inner Join



b

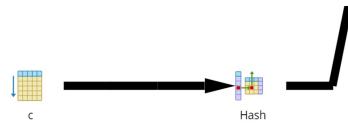


Hash



Hash Inner Join

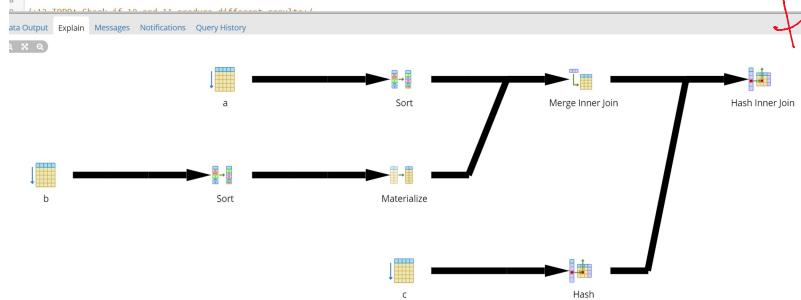
- limited



$T(A) = 5000000$ $V(A, ht) = 100000$
 $T(B) = 5000000$ $V(B, ht) = 100000$
 $T(c) = 5000000$ $V(C, ht) = 100000$
 Selected plan: hash join A,b, hash join c

11. A.ten = B.ten = C.ten

```
5
6
7 /*11*/
```



 Optimized
has decided it's
better to use hashing
than to sort the 10^5
values in the attribute.

T(A) = 5000000 V(A,ten) = 10
T(B) = 5000000 V(B,ten) = 10
T(c) = 5000000 V(C,ten) = 10
Selected plan: merge join A,B, hash join c

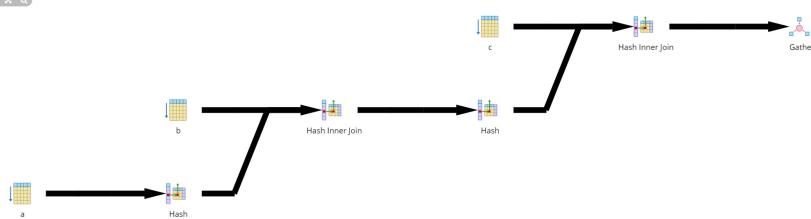
12. Do they produce different plans?

Yes, the plans are different. I found the plan changes on the ot attribute when $V(R,ot) \neq 1000$
 $A_ot = B_ot = C_ot$

A.ot = B.ot = C.ot

13. A.pk = B.ht = C.ht

```
12 SELECT * FROM A JOIN B ON A.a1 = B.b1 JOIN C ON B.b1 = C.c1; */
13
14 /*13*/
15 SELECT * FROM A JOIN B ON A.apk = B.ht JOIN C ON B.ht = C.ht;
16
17 /*14*/
18 SELECT * FROM A JOIN B ON A.apk = B.ht JOIN C ON B.ht = C.ht;
```



A optimizer has decided it's better to sort the 10 values in the table and merge join.

Selected plan is to hash join a,b then hash join c
T(A) = 5000000 V(A,pk) = 5000000
T(B) = 5000000 V(B,ht) = 100000
T(C) = 5000000 V(C,ht) = 100000

14. A.pk = B.hund = C.hund

```
SELECT * FROM A JOIN B ON A.pk = B.hund JOIN C ON B.hund = C.hund
```

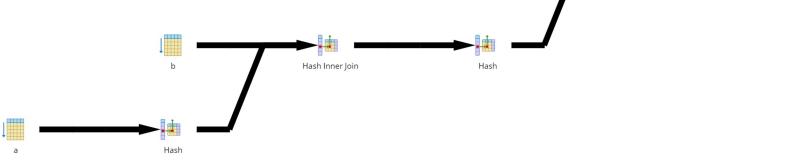
```
/*15 TODO: Check if 13 and 14 produce different results */
```

```
/*16*/  
SELECT * FROM A_prime JOIN B_prime ON A_prime
```

/*17*/

a Output Explain Messages Notifications Query History

2023-01-10



Selected plan is to hash join a,b then hash join c

$$T(A) = 5000000 \quad V(A, \text{pk}) = 5000000$$

$$T(B) = 5000000 \quad V(B, \text{hund}) = 100$$

$$T(C) = 5000000 \quad V(C, \text{hund}) = 100$$

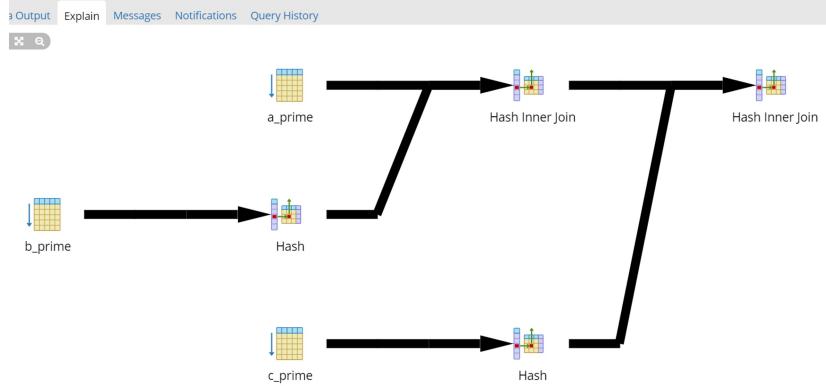
15. Do they produce different plans?

No, 13 and 14 produce the same plan.

16. $A'.ht = B'.ht = C'.ht$

```
/*16*/
SELECT * FROM A_prime JOIN B_prime ON A_prime.ht = B_prime.ht JOIN C_prime ON B_prime.ht = C_prime.ht

/*17*/
SELECT * FROM A_prime JOIN B_prime ON A_prime.hund = B_prime.hund JOIN C_prime ON B_prime.hund = C_prime.hund
```



$$T(A) = 5000000 \quad V(A, ht) = 100000$$

$$T(B) = 5000000 \quad V(B, ht) = 100000$$

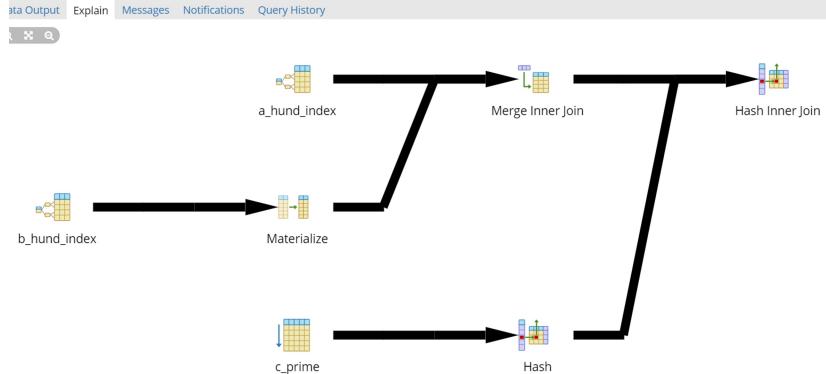
$$T(C) = 5000000 \quad V(C, ht) = 100000$$

Selected plan: hash join A,b, hash join c. The query plan does not change from 10, so the secondary index has no effect on the plan

17. $A'.hund = B'.hund = C'.hund$

```
/*17*/
SELECT * FROM A_prime JOIN B_prime ON A_prime.hund = B_prime.hund JOIN C_prime ON B_prime.hund = C_prime.hund

/*18*/
SELECT * FROM A_prime JOIN B_prime ON A_prime.ht = B_prime.ten JOIN C_prime ON B_prime.ten = C_prime.ht
```



$$T(A) = 5000000 \quad V(A, hund) = 100$$

$$T(B) = 5000000 \quad V(B, hund) = 100$$

$$T(C) = 5000000 \quad V(C, hund) = 100$$

Selected plan: merge join A,b, hash join c. The query plan changes from 16.

18. Do 16 and 17 produce different results?

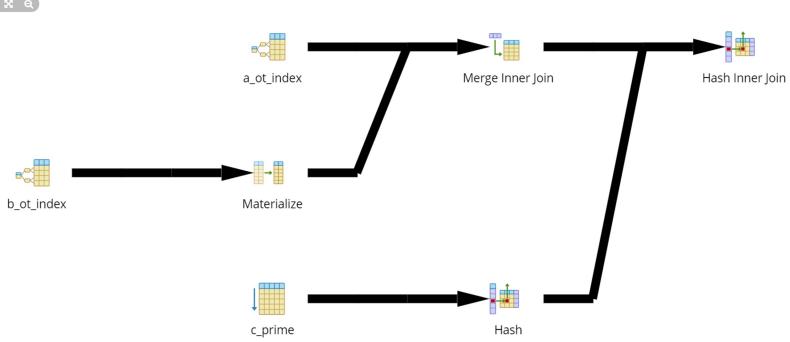
Yes. 16 chooses hash join to join a,b. 17 chooses merge join. The transition point is at the ht attribute

```

3 /*Transition point*/
4 SELECT * FROM A_prime JOIN B_prime ON A_prime.ot = B_prime.ot JOIN C_prime ON B_prime.ot = C_prime.ot
5
6 /*19 */
7 SELECT * FROM A_prime,C_prime,A,B,B_prime,C
8 WHERE
9 A_prime.ht = C_prime.ot AND
10 A.ht = A_prime.ten AND
11 B.pk = C_prime.hund AND
12

```

Data Output Explain Messages Notifications Query History



Section 2:

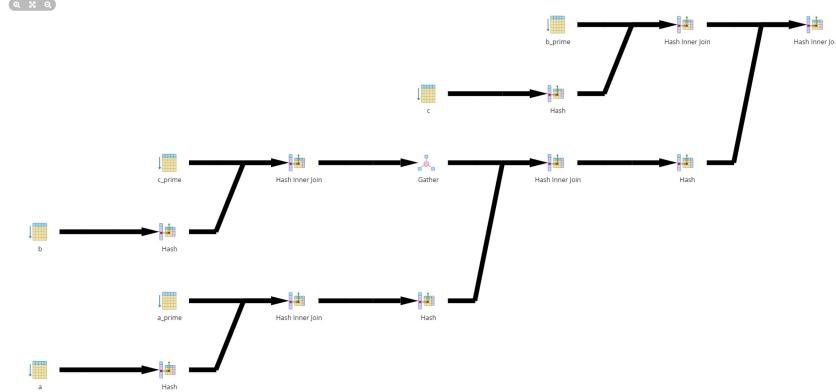
19. $A'.ht = C'.ot \text{ AND}$
 $A.ht = A'.ten \text{ AND}$
 $B.pk = C'.hund \text{ AND}$
 $A'.ht = C'.ot \text{ AND}$
 $B'.ten = C.ot \text{ AND}$
 $B'.hund = A.ht$

```

102 /*19 */
104 SELECT * FROM A_prime,C_prime,A,B,B_prime,C
105 WHERE
106 A_prime.ht = C_prime.ot AND
107 A.ht = A_prime.ten AND
108 B.pk = C_prime.hund AND
109 A_prime.ht = C_prime.ot AND
110 B_prime.ten = C.ot AND
111 B_prime.hund = A.ht
112

```

Data Output Explain Messages Notifications Query History



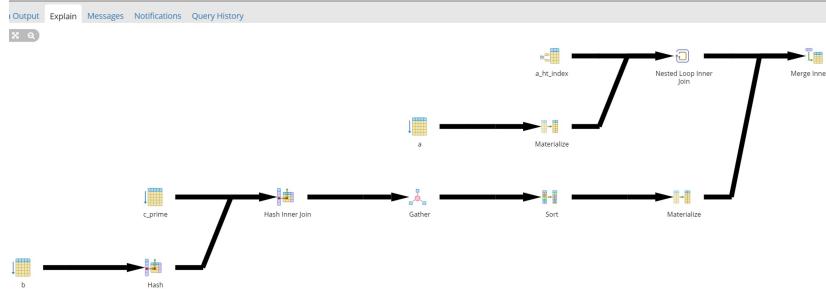
The estimated size in the last hash inner join is over 3 quadrillion rows (3,172,245,629,574,755)

20. $A'.ht = C'.ot \text{ AND}$
 $A'.ten = 5 \text{ AND}$
 $A.ht < A'.ten \text{ AND}$
 $B.pk = C'.hund \text{ AND}$
 $B.ot < 500 \text{ AND}$
 $A'.ht = C'.ot$

```

/*+20*/
SELECT * FROM A_prime,C_prime,A,B
WHERE
A_prime.ht = C_prime.ot AND
A_prime.tt = 5 AND
A_prime.A_prime_ten AND
B_pk < C_prime.hund AND
B_tt < 500 AND
A_prime.ht = C_prime.ot

```



The estimated size of the last merge inner join is: over 21 trillion rows(21,554,197,796,542)!

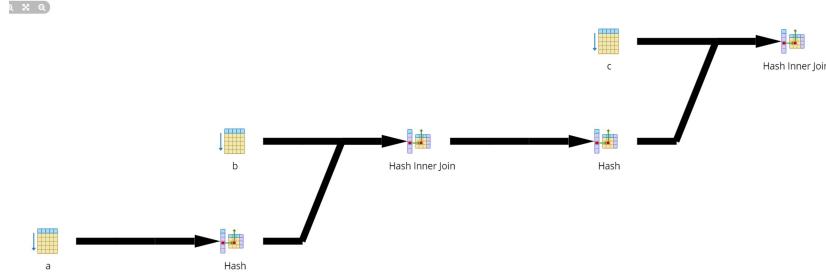
The compiler did not notice the equivalent predicate.

21. A.ht = B.tt = C.ot

```

13 /*+21*/
14 SELECT * FROM A
15 JOIN B ON A.ht = B.tt
16 JOIN C ON B.tt = C.ot
17
18 /*+22*/
19 SELECT * FROM C
20 JOIN B ON C.ot = B.tt

```



$$T(A) = 5000000 \quad V(A,ht) = 100000$$

$$T(B) = 5000000 \quad V(B,tt) = 10000$$

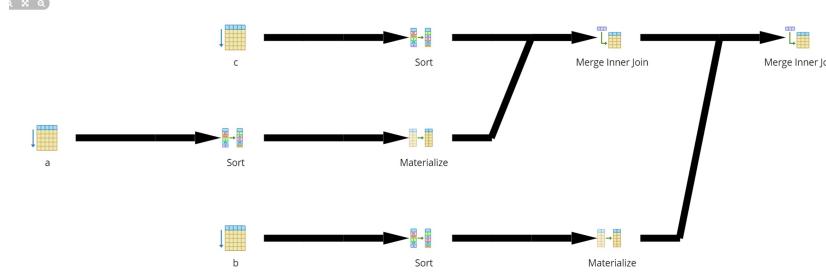
$$T(C) = 5000000 \quad V(C,ot) = 1000$$

22. C.ot = B.tt = A.ht

```

7
8 /*+22*/
9 SELECT * FROM C
10 JOIN B ON C.ot = B.tt
11 JOIN A ON B.tt = A.ht
12

```



The query plan changes between 21 and 22 even though the only difference is the order of the tables in the query.

$$T(A) = 5000000 \quad V(A,ht) = 100000$$

$$T(B) = 5000000 \quad V(B,tt) = 10000$$

$$T(C) = 5000000 \quad V(C,ot) = 1000$$

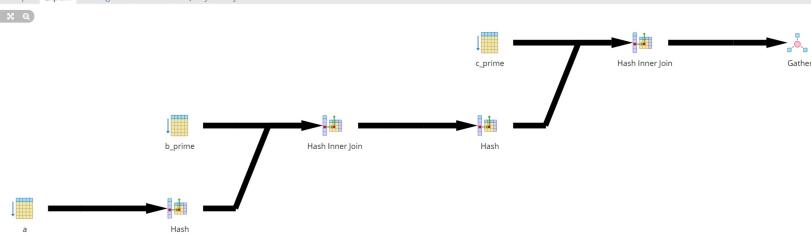
23. A.pk = B'.tt = C'.ot

```

32  /*+2x*/
33  SELECT * FROM A
34  JOIN B_prime ON A.pk = B_prime.tt
35  JOIN C_prime ON B_prime.tt = C_prime.ot
36
37  /*+2x*/
38  SELECT *FROM C_prime
39  JOIN B_prime ON C_prime.ot = B_prime.tt
40  JOIN A ON B_prime.tt = A.pk
41
42

```

Data Output Explain Messages Notifications Query History



$$T(A) = 5000000 \quad T(A,pk) = 5000000$$

$$T(B) = 5000000 \quad T(B,tt) = 10000$$

$$T(C) = 5000000 \quad T(C,ot) = 1000$$

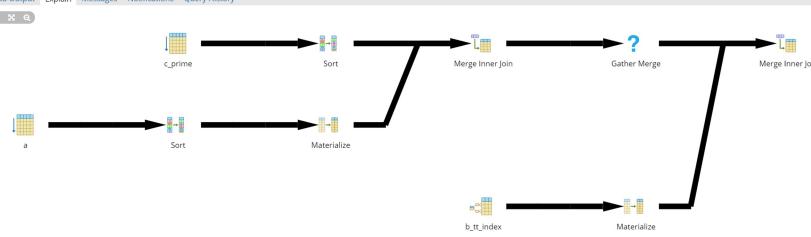
24. $C'.ot = B'.tt = A.pk$

```

37  /*+2x*/
38  SELECT *FROM C_prime
39  JOIN B_prime ON C_prime.ot = B_prime.tt
40  JOIN A ON B_prime.tt = A.pk
41
42

```

Data Output Explain Messages Notifications Query History



$$T(A) = 5000000 \quad T(A,pk) = 5000000$$

$$T(B) = 5000000 \quad T(B,tt) = 10000$$

$$T(C) = 5000000 \quad T(C,ot) = 1000$$

Again, the query plan changes between 23 and 24 even though the only difference is the order of the tables in the query